

## Git Reference

- [Reference](#)
- [About](#)
- §
- [Site Source](#)

## Getting and Creating Projects

- [init](#)
- [clone](#)

## Basic Snapshotting

- [add](#)
- [status](#)
- [diff](#)
- [commit](#)
- [reset](#)
- [rm, mv](#)
- [stash](#)

## Branching and Merging

- [branch](#)
- [checkout](#)
- [merge](#)
- [log](#)

- [tag](#)

## [Sharing and Updating Projects](#)

- [remote](#)
- [fetch, pull](#)
- [push](#)

## [Inspection and Comparison](#)

- [log](#)
- [diff](#)

## [book](#) Basic Snapshotting

Git is all about composing and saving snapshots of your project and then working with and comparing those snapshots. This section will explain the commands needed to compose and commit snapshots of your project.

An important concept here is that Git has an 'index', which acts as sort of a staging area for your snapshot. This allows you to build up a series of well composed snapshots from changed files in your working directory, rather than having to commit all of the file changes at once.

**In a nutshell**, you will use `git add` to start tracking new files and also to stage changes to already tracked files, then `git status` and `git diff` to see what has been modified and staged and finally `git commit` to record your snapshot into your history. This will be the basic workflow that you use most of the time.

## [docs](#) [book](#) git add adds file contents to the staging area

In Git, you have to add file contents to your staging area before you can commit them. If the file is new, you can run `git add` to initially add the file to your staging area, but even if the file is already "tracked" - ie, it was in your last commit - you still need to call `git add` to add new modifications to your staging area. Let's see a

few examples of this.

Going back to our Hello World example, once we've initiated the project, we would now start adding our files to it and we would do that with `git add`. We can use `git status` to see what the state of our project is.

```
$ git status -s
?? README
?? hello.rb
```

So right now we have two untracked files. We can now add them.

```
$ git add README hello.rb
```

Now if we run `git status` again, we'll see that they've been added.

```
$ git status -s
A README
A hello.rb
```

It is also common to recursively add all files in a new project by specifying the current working directory like this: `git add ..`. Since Git will recursively add all files under a directory you give it, if you give it the current working directory, it will simply start tracking every file there. In this case, a `git add .` would have done the same thing as a `git add README hello.rb`, or for that matter so would `git add *`, but that's only because we don't have subdirectories which the `*` would not recurse into.

OK, so now if we edit one of these files and run `git status` again, we will see something odd.

```
$ vim README
$ git status -s
AM README
A hello.rb
```

The 'AM' status means that the file has been modified on disk since we last added it. This means that if we commit our snapshot right now, we will be recording the version of the file when we last ran `git add`, not the version that is on our disk. Git does not assume that what the file looks like on disk is necessarily what you want to snapshot - you have to tell Git with the `git add` command.

**In a nutshell**, you run `git add` on a file when you want to include whatever changes you've made to it in your next commit snapshot. Anything you've changed that is not added will not be included - this means you can craft your snapshots with a bit more precision than most other SCM systems.

For a very interesting example of using this flexibility to stage only parts of modified files at a time, see the '-p' option to `git add` in the Pro Git book.

## [docs](#) [book](#) **git status view the status of your files in the working directory and staging area**

As you saw in the `git add` section, in order to see what the status of your staging area is compared to the code in your working directory, you can run the `git status` command. Using the `-s` option will give you short output. Without that flag, the `git status` command will give you more context and hints. Here is the same status output with and without the `-s`. The short output looks like this:

```
$ git status -s
AM README
A  hello.rb
```

Where the same status with the long output looks like this:

```
$ git status
# On branch master
#
# Initial commit
#
# Changes to be committed:
#   (use "git rm --cached <file>..." to unstage)
#
# new file:   README
# new file:   hello.rb
#
# Changed but not updated:
#   (use "git add <file>..." to update what will be committed)
#   (use "git checkout -- <file>..." to discard changes in working directory)
```

```
#  
# modified:  README  
#
```

You can easily see how much more compact the short output is, but the long output has useful tips and hints as to what commands you may want to use next.

Git will also tell you about files that were deleted since your last commit or files that were modified or staged since your last commit.

```
$ git status -s  
M README  
D hello.rb
```

You can see there are two columns in the short status output. The first column is for the staging area, the second is for the working directory. So for example, if you have the README file staged and then you modify it again without running `git add` a second time, you'll see this:

```
$ git status -s  
MM README  
D hello.rb
```

**In a nutshell**, you run `git status` to see if anything has been modified and/or staged since your last commit so you can decide if you want to commit a new snapshot and what will be recorded in it.

## [docs](#) [book](#) **git diff shows diff of what is staged and what is modified but unstaged**

There are two main uses of the `git diff` command. One use we will describe here, the other we will describe later in the ["Inspection and Comparison"](#) section. The way we're going to use it here is to describe the changes that are staged or modified on disk but unstaged.

### **git diff show diff of unstaged changes**

Without any extra arguments, a simple `git diff` will display in unified diff format (a patch) what code or content you've changed in your project since the last commit that are not yet staged for the next commit snapshot.

```
$ vim hello.rb
$ git status -s
 M hello.rb
$ git diff
diff --git a/hello.rb b/hello.rb
index d62ac43..8d15d50 100644
--- a/hello.rb
+++ b/hello.rb
@@ -1,7 +1,7 @@
 class HelloWorld

   def self.hello
-    puts "hello world"
+    puts "hola mundo"
   end

 end
```

So where `git status` will show you what files have changed and/or been staged since your last commit, `git diff` will show you what those changes actually are, line by line. It's generally a good follow-up command to `git status`

### **`git diff --cached` show diff of staged changes**

The `git diff --cached` command will show you what contents have been staged. That is, this will show you the changes that will currently go into the next commit snapshot. So, if you were to stage the change to `hello.rb` in the example above, `git diff` by itself won't show you any output because it will only show you what is *not yet* staged.

```
$ git status -s
 M hello.rb
$ git add hello.rb
```

```
$ git status -s
M hello.rb
$ git diff
$
```

If you want to see the staged changes, you can run `git diff --cached` instead.

```
$ git status -s
M hello.rb
$ git diff
$
$ git diff --cached
diff --git a/hello.rb b/hello.rb
index d62ac43..8d15d50 100644
--- a/hello.rb
+++ b/hello.rb
@@ -1,7 +1,7 @@
  class HelloWorld

    def self.hello
-    puts "hello world"
+    puts "hola mundo"
    end

  end
```

### **git diff HEAD** show diff of all staged or unstaged changes

If you want to see both staged and unstaged changes together, you can run `git diff HEAD` - this basically means you want to see the difference between your working directory and the last commit, ignoring the staging area. If we make another change to our `hello.rb` file then we'll have some changes staged and some changes unstaged. Here are what all three `diff` commands will show you:

```
$ vim hello.rb
$ git diff
diff --git a/hello.rb b/hello.rb
index 4f40006..2ae9ba4 100644
```

```
--- a/hello.rb
+++ b/hello.rb
@@ -1,7 +1,7 @@
  class HelloWorld

+   # says hello
    def self.hello
      puts "hola mundo"
    end

  end
$ git diff --cached
diff --git a/hello.rb b/hello.rb
index 2aabb6e..4f40006 100644
--- a/hello.rb
+++ b/hello.rb
@@ -1,7 +1,7 @@
  class HelloWorld

    def self.hello
-     puts "hello world"
+     puts "hola mundo"
    end

  end
$ git diff HEAD
diff --git a/hello.rb b/hello.rb
index 2aabb6e..2ae9ba4 100644
--- a/hello.rb
+++ b/hello.rb
@@ -1,7 +1,8 @@
  class HelloWorld

+   # says hello
    def self.hello
-     puts "hello world"
+     puts "hola mundo"
    end

  end
```



## **git diff --stat** show summary of changes instead of a full diff

If we don't want the full diff output, but we want more than the `git status` output, we can use the `--stat` option, which will give us a summary of changes instead. Here is the same example as above, but using the `--stat` option instead.

```
$ git status -s
MM hello.rb
$ git diff --stat
hello.rb | 1 +
1 files changed, 1 insertions(+), 0 deletions(-)
$ git diff --cached --stat
hello.rb | 2 +-
1 files changed, 1 insertions(+), 1 deletions(-)
$ git diff HEAD --stat
hello.rb | 3 ++-
1 files changed, 2 insertions(+), 1 deletions(-)
```

You can also provide a file path at the end of any of these options to limit the `diff` output to a specific file or subdirectory.

**In a nutshell**, you run `git diff` to see details of the `git status` command - *how* files have been modified or staged on a line by line basis.

## **[docs](#) [book](#) git commit records a snapshot of the staging area**

Now that you have staged the content you want to snapshot with the `git add` command, you run `git commit` to actually record the snapshot. Git records your name and email address with every commit you make, so the first step is to tell Git what these are.

```
$ git config --global user.name 'Your Name'
$ git config --global user.email you@somedomain.com
```

Let's stage and commit all the changes to our `hello.rb` file. In this first example, we'll use the `-m` option to provide the commit message on the command line.

```
$ git add hello.rb
$ git status -s
M hello.rb
$ git commit -m 'my hola mundo changes'
[master 68aa034] my hola mundo changes
1 files changed, 2 insertions(+), 1 deletions(-)
```

Now we have recorded the snapshot. If we run `git status` again, we will see that we have a "clean working directory", which means that we have not made any changes since our last commit - there is no un-snapshotted work in our checkout.

```
$ git status
# On branch master
nothing to commit (working directory clean)
```

If you leave off the `-m` option, Git will try to open a text editor for you to write your commit message. In `vim`, which it will default to if it can find nothing else in your settings, the screen might look something like this:

```
# Please enter the commit message for your changes. Lines starting
# with '#' will be ignored, and an empty message aborts the commit.
# On branch master
# Changes to be committed:
#   (use "git reset HEAD <file>..." to unstage)
#
# modified:   hello.rb
#
~
~
".git/COMMIT_EDITMSG" 9L, 257C
```

At this point you add your actual commit message at the top of the document. Any lines starting with `'#'` will be ignored - Git will put the output of the `git status` command in there for you as a reminder of what you have modified and staged.

In general, it's very important to write a good commit message. For open source projects, it's generally a rule to write your message more or less in this format:

Short (50 chars or less) summary of changes

More detailed explanatory text, if necessary. Wrap it to about 72 characters or so. In some contexts, the first line is treated as the subject of an email and the rest of the text as the body. The blank line separating the summary from the body is critical (unless you omit the body entirely); some git tools can get confused if you run the two together.

Further paragraphs come after blank lines.

- Bullet points are okay, too
- Typically a hyphen or asterisk is used for the bullet, preceded by a single space, with blank lines in between, but conventions vary here

```
# Please enter the commit message for your changes. Lines starting
# with '#' will be ignored, and an empty message aborts the commit.
# On branch master
# Changes to be committed:
#   (use "git reset HEAD <file>..." to unstage)
#
# modified:   hello.rb
#
~
~
~
".git/COMMIT_EDITMSG" 25L, 884C written
```

The commit message is very important. Since much of the power of Git is this flexibility in carefully crafting commits locally and then sharing them later, it is very powerful to be able to write three or four commits of logically separate changes so that your work may be more easily peer reviewed. Since there is a separation between committing and pushing those changes, do take the time to make it easier for the people you are working with to see what you've done by putting each logically separate change in a separate commit with a nice commit message so it is easier for them to see what you are doing and why.

**git commit -a** automatically stage all tracked, modified files before the commit

If you think the `git add` stage of the workflow is too cumbersome, Git allows you to skip that part with the `-a` option. This basically tells Git to run `git add` on any file that is "tracked" - that is, any file that was in your last commit and has been modified. This allows you to do a more Subversion style workflow if you want, simply editing files and then running `git commit -a` when you want to snapshot everything that has been changed. You still need to run `git add` to start tracking new files, though, just like Subversion.

```
$ vim hello.rb
$ git status -s
M hello.rb
$ git commit -m 'changes to hello file'
# On branch master
# Changed but not updated:
#   (use "git add <file>..." to update what will be committed)
#   (use "git checkout -- <file>..." to discard changes in working directory)
#
# modified:   hello.rb
#
no changes added to commit (use "git add" and/or "git commit -a")
$ git commit -am 'changes to hello file'
[master 78b2670] changes to hello file
1 files changed, 2 insertions(+), 1 deletions(-)
```

Notice how if you don't stage any changes and then run `git commit`, Git will simply give you the output of the `git status` command, reminding you that nothing is staged. The important part of that message has been highlighted, saying that nothing is added to be committed. If you use `-a`, it will add and commit everything at once.

This now lets you complete the entire snapshotting workflow - you make changes to your files, then use `git add` to stage files you want to change, `git status` and `git diff` to see what you've changed, and then finally `git commit` to actually record the snapshot forever.

**In a nutshell**, you run `git commit` to record the snapshot of your staged content. This snapshot can then be compared, shared and reverted to if you need to.

[docs](#) [book](#) **git reset undo changes and commits**

`git reset` is probably the most confusing command written by humans, but it can be very useful once you get the hang of it. There are three specific invocations of it that are generally helpful.

### **`git reset HEAD` unstage files from index and reset pointer to HEAD**

First, you can use it to unstage something that has been accidentally staged. Let's say that you have modified two files and want to record them into two different commits. You should stage and commit one, then stage and commit the other. If you accidentally stage both of them, how do you *un-stage* one? You do it with `git reset HEAD -- file`. Technically you don't have to add the `--` - it is used to tell Git when you have stopped listing options and are now listing file paths, but it's probably good to get into the habit of using it to separate options from paths even if you don't need to.

Let's see what it looks like to unstage something. Here we have two files that have been modified since our last commit. We will stage both, then unstage one of them.

```
$ git status -s
M README
M hello.rb
$ git add .
$ git status -s
M README
M hello.rb
$ git reset HEAD -- hello.rb
Unstaged changes after reset:
M hello.rb
$ git status -s
M README
M hello.rb
```

Now you can run a `git commit` which will just record the changes to the README file, not the now unstaged hello.rb.

In case you're curious, what it's actually doing here is it is resetting the checksum of the entry for that file in the "index" to be what it was in the last commit. Since `git add` checksums a file and adds it to the "index", `git reset HEAD` overwrites that with what it was before, thereby effectively unstaging it.

If you want to be able to just run `git unstage`, you can easily setup an alias in Git. Just run `git config --global alias.unstage "reset HEAD"`. Once you have run that, you can then just run `git unstage [file]` instead.

If you forget the command to unstage something, Git is helpful in reminding you in the output of the normal `git status` command. For example, if you run `git status` without the `-s` when you have staged files, it will tell you how to unstage them:

```
$ git status
# On branch master
# Changes to be committed:
#   (use "git reset HEAD <file>..." to unstage)
#
#   modified:   README
#   modified:   hello.rb
#
```

When you run `git reset` without specifying a flag it defaults to `--mixed`. The other options are `--soft` and `--hard`.

### **git reset --soft moves HEAD to specified commit reference, index and staging are untouched**

The first thing `git reset` does is undo the last commit and put the files back onto the stage. If you include the `--soft` flag this is where it stops. For example, if you run `git reset --soft HEAD~` (the parent of the HEAD) the last commit will be undone and the files touched will be back on the stage again.

```
$ git status -s
M hello.rb
$ git commit -am 'hello with a flower'
[master 5857acl] hello with a flower
1 files changed, 3 insertions(+), 1 deletions(-)
$ git status
# On branch master
nothing to commit (working directory clean)
$ git reset --soft HEAD~
$ git status -s
M hello.rb
```

This is basically doing the same thing as `git commit --amend`, allowing you to do more work before you roll in the file changes into the same commit.

### **git reset --hard unstage files AND undo any changes in the working directory since last commit**

The third option is to go `--hard`. This command discards your staged changes and the changes in your working directory. In other words: it resets your staging area and working directory to the state they were in at the given commit. This is the most dangerous option and is not working directory safe. Any changes not committed will be lost.

```
$ git status
# On branch master
# Changes to be committed:
#   (use "git reset HEAD <file>..." to unstage)
#
# modified:   README
#
# Changes not staged for commit:
#   (use "git add <file>..." to update what will be committed)
#   (use "git checkout -- <file>..." to discard changes in working directory)
#
# modified:   README
#
$ git reset --hard HEAD
HEAD is now at 5857ac1 hello with a flower
$ git status
# On branch master
nothing to commit (working directory clean)
```

In the above example, while we had both changes ready to commit and ready to stage, a `git reset --hard` wiped them out. The working tree and staging area are reset to the tip of the current branch or HEAD.

You can replace HEAD with a commit SHA-1 or another parent reference to reset to that specific point.

**In a nutshell**, you run `git reset HEAD` to undo the last commit, unstage files that you previously ran `git add` on and wish to not include in the next commit snapshot

## [docs](#) [book](#) **git rm remove files from the staging area**

`git rm` will remove entries from the staging area. This is a bit different from `git reset HEAD` which "unstages" files. To "unstage" means it reverts the staging area to what was there before we started modifying things. `git rm` on the other hand just kicks the file off the stage entirely, so that it's not included in the next commit snapshot, thereby effectively deleting it.

By default, a `git rm file` will remove the file from the staging area entirely and also off your disk (the working directory). To leave the file in the working directory, you can use `git rm --cached .`

**git mv** `git rm --cached orig; mv orig new; git add new`

Unlike most other version control systems, Git does not track file renames. Instead, it just tracks the snapshots and then figures out what files were likely renamed by comparing snapshots. If a file was removed from one snapshot and another file was added to the next one and the contents are similar, Git figures it was most likely a rename. So, although the `git mv` command exists, it is superfluous - all it does is a `git rm --cached`, moves the file on disk, then runs a `git add` on the new file. You don't really need to use it, but if it's easier, feel free.

In its normal form the command is used to delete files. But it's often easier to just remove the files off your disk and then run `git commit -a`, which will also automatically remove them from your index.

**In a nutshell**, you run `git rm` to remove files from being tracked in Git. It will also remove them from your working directory.

## [docs](#) [book](#) **git stash save changes made in the current index and working directory for later**

You're in the middle of some changes but something comes up that you need to jump over to, like a so-urgent-right-now bugfix, but don't want to commit or lose your current edits. `git stash` is there for you.



## **git stash add current changes to the stack**

Stashing takes the current state of the working directory and index, puts it on a stack for later, and gives you back a clean working directory. It will then leave you at the state of the last commit.

If you have untracked files, `git stash` will not include them. You can either stage those files with `git add` (you don't have to commit) before stashing, or, if you have a recent Git version (1.7.7 or above), you can use `git stash -u` to also stash also unversioned files.

```
$ git status -s
M hello.rb
$ git stash
Saved working directory and index state WIP on master: 5857ac1 hello with a flower
HEAD is now at 5857ac1 hello with a flower
$ git status
# On branch master
nothing to commit (working directory clean)
```

## **git stash list view stashes currently on the stack**

It's helpful to know what you've got stowed on the stash and this is where `git stash list` comes in. Running this command will display a queue of current stash items.

```
$ git stash list
stash@{0}: WIP on master: 5857ac1 hello with a flower
```

The last item added onto the stash will be referenced by `stash@{0}` and increment those already there by one.

```
$ vim hello.rb
$ git commit -am 'it stops raining'
[master ee2d2c6] it stops raining
1 files changed, 1 insertions(+), 1 deletions(-)
$ vim hello.rb
$ git stash
Saved working directory and index state WIP on master: ee2d2c6 it stops raining
HEAD is now at ee2d2c6 it stops raining
```

```
$ git stash list
stash@{0}: WIP on master: ee2d2c6 it stops raining
stash@{1}: WIP on master: 5857ac1 hello with a flower
```

### **git stash apply** grab the item from the stash list and apply to current working directory

When you're ready to continue from where you left off, run the `git stash apply` command to bring back the saved changes onto the working directory.

```
$ git stash apply
# On branch master
# Changes not staged for commit:
#   (use "git add <file>..." to update what will be committed)
#   (use "git checkout -- <file>..." to discard changes in working directory)
#
# modified:   hello.rb
#
no changes added to commit (use "git add" and/or "git commit -a")
```

By default it will reapply the last added stash item to the working directory. This will be the item referenced by `stash@{0}`. You can grab another stash item instead if you reference it in the arguments list. For example, `git stash apply stash@{1}` will apply the item referenced by `stash@{1}`.

If you also want to remove the item from the stack at the same time, use `git stash pop` instead.

### **git stash drop** remove an item from the stash list

When you're done with the stashed item and/or want to remove it from the list, run the `git stash drop` command. By default this will remove the last added stash item. You can also remove a specific item if you include it as an argument.

In this example, our stash list has at least two items, but we want to get rid of the item added before last, which is referenced by `stash@{1}`.

```
$ git stash drop stash@{1}
```

Dropped stash@{1} (0b1478540189f30fef9804684673907c65865d8f)

If you want to remove of all the stored items, just run the `git stash clear` command. But only do this if you're sure you're done with the stash.

**In a nutshell**, run `git stash` to quickly save some changes that you're not ready to commit or save, but want to come back to while you work on something else.

[On to Branching and Merging »](#)